

DOM-Graph RAG: DOM-Guided Adaptive Tiling and Budgeted Online Evidence Control for Multimodal Web-Page RAG

Anonymous ACL submission

Abstract

Recent visual retrieval-augmented generation (RAG) systems render documents as pixel screenshots and let a vision-language model (VLM) read tiles of the image directly, avoiding the information loss of OCR and text extraction. PixelRAG (Wang et al., 2026) popularized this idea but tiles the rendered page on a fixed pixel grid, blind to content boundaries: a table or paragraph can be sliced in half mid-tile. Separately, MAGE-RAG (Zuo et al., 2026) organizes long documents into a multigranular page/element evidence graph and reads it with an online controller that activates, opens, and prunes graph nodes under a fixed action budget, but targets already-parsed long PDFs rather than rendered web pages. We present **DOM-Graph RAG**, a pipeline that replaces PixelRAG’s blind grid with tiling guided by the page’s own DOM structure, so tile boundaries follow content boundaries by construction. The resulting tiles are organized into a MAGE-RAG-style multigranular graph, extended across pages via hyperlink edges into a small corpus graph. We fine-tune a VLM reader contrastively on synthetically generated (tile, question) pairs with mined hard negatives (LoRA, both vision tower and decoder, InfoNCE), and implement an online evidence controller that evolves every node through a four-state machine (INACTIVE→ACTIVE→OPENED/PRUNED) via best-first search under budget, so the reader gets a small, prioritized evidence sub-graph instead of the whole page. We describe the full, open-sourced pipeline in detail and report a first preliminary result: on a slice of our own synthetic QA set, DOM-guided tiling’s lexical retrieval signal is 14–16× stronger than chance, against 3.7–3.9× for a blind fixed-height grid at the same resolution, once pool-size differences between the two tilings are controlled for. Full benchmark evaluation against MAGE-RAG’s published LongDocURL and MMLongBench-Doc

numbers is ongoing work, reported as such in Section 5.

1 Introduction

Retrieval-augmented generation (RAG) over visually rich documents – web pages, tables, infoboxes – has long relied on text extraction: parse the HTML or run OCR, chunk the text, retrieve with a dense retriever. This discards exactly what layout encodes: which numbers belong to which table row, which caption belongs to which figure. PixelRAG (Wang et al., 2026) instead shows that a VLM retrieving directly over rendered page screenshots can match or beat text-based retrieval on long, layout-heavy documents – but its tiling is a fixed pixel grid, blind to what is on the page: every tile has the same height regardless of content, so a table row or paragraph can be split with no awareness it was ever one unit. A retriever built on such tiles must reconstruct boundaries the renderer already knew and threw away.

MAGE-RAG (Zuo et al., 2026) addresses a different problem: given an already-segmented long document, which passages should an agent read under a bounded budget? It builds a multigranular evidence graph (containment, reading order, layout adjacency, section hierarchy, semantic neighbors) and an online controller (its Algorithm 1) that activates, opens, and prunes nodes under budget – but assumes the document is already parsed, and does not address obtaining content-aligned tiles from a rendered page in the first place.

The two are complementary: PixelRAG shows pixel tiles read by a VLM are a viable retrieval unit, but tiles blindly; MAGE-RAG shows a budget-aware graph controller beats static top- k , but assumes elements already exist. This paper combines them via the one artifact PixelRAG discards and MAGE-RAG never had to construct: the DOM. For a rendered web page, the browser’s own layout engine already knows every

element’s exact bounding box before a single pixel is cropped.

Contributions. We introduce **DOM-Graph RAG**, an open-source pipeline (§3) that: (1) replaces PixelRAG’s blind grid with *DOM-guided adaptive tiling* – elements cropped along their own bounding boxes, undersized ones merged, oversized ones split, headings always paired with their content, nested duplicates pruned before tiling (§3.1); (2) organizes tiles into a *multigranular graph* in the style of MAGE-RAG – all five of its relations, plus cross-page hyperlink edges into a small corpus graph, and oversized pages split into section-level sub-pages via a new *continues* relation, so document length grows the graph rather than excluding long documents (§3.2); (3) fine-tunes a general-purpose VLM into a page-element reader via *contrastive LoRA training* on synthetic (tile, question) pairs with lexically- and, optionally, VLM-judge-mined hard negatives, mirrored at reduced resolution for a controlled test of PixelRAG’s resolution/token-cost claim (§3.3); (4) implements an *online evidence controller* generalizing MAGE-RAG’s Algorithm 1 to live-DOM graphs via an explicit four-state machine under a node-opening budget (§3.4).

The full pipeline is implemented and exercised end-to-end on live Wikipedia pages, with dependency-light unit tests over the geometric, graph, mining, and image-processing logic. §5 reports a first preliminary ablation and lays out the evaluation against MAGE-RAG’s published baselines as work in progress, rather than fabricate numbers ahead of running it.

2 Related Work

Text-based document RAG. Standard RAG chunks extracted text and retrieves with a dense bi-encoder (Karpukhin et al., 2020). On visually rich documents this requires a layout-aware parser or OCR, both lossy on complex layouts, and discards 2D structure (which text belongs to which cell or caption) before retrieval ever runs.

Screenshot-based visual RAG. PixelRAG (Wang et al., 2026) retrieves over rendered-page image tiles with a VLM embedding model, part of an established line retrieving over page images rather than text: ColPali (Faysse et al., 2024) (late-interaction patch embeddings), Doc-

ument Screenshot Embedding (Ma et al., 2024a) (whole-page VLM embeddings, no parsing), and VisRAG (Yu et al., 2024)/M3DocRAG (Cho et al., 2024) (full retrieve-then-generate pipelines, the latter multi-page/-document). None address *how* a page is cut into retrievable units – what our tiling (§3.1) targets, so it could feed any of them, not only PixelRAG, whose own blind fixed-height grid lets a table or paragraph straddle a tile boundary. PixelRAG also reports up to $3\times$ token-cost reduction at lower resolution without specifying the filter or resolutions tested; we build a matched full-/reduced-resolution pair on our own tiles to measure this directly (§3.3).

Graph-structured evidence and controllers. MAGE-RAG (Zuo et al., 2026) (evaluated on LongDocURL (Deng et al., 2025), MMLongBench-Doc (Ma et al., 2024b)) builds an offline multigranular evidence graph and reads it with a budgeted controller – the direct ancestor of §3.4; the difference is provenance: we build the graph from a rendered page’s DOM, giving exact boundaries and reading order for free, and now instantiate all five of its relation types (§3.2). GraphRAG (Edge et al., 2024) shares the premise of building an explicit graph before querying it, over extracted entities rather than layout structure – complementary rather than competing. Layout- and markup-aware pretraining (LayoutLM (Xu et al., 2020), MarkupLM (Li et al., 2022)) instead bakes structure into the model’s representation; we use the DOM directly at inference time, with no pretraining required. Web-agent work (WebArena (Zhou et al., 2023), Mind2Web (Deng et al., 2023), SeeAct (Zheng et al., 2024)) has already studied when DOM structure is reliable enough to act on – the same failure mode (JS-hydrated, div-soup markup) that scopes our own claim to semantically marked-up pages (§6).

Contrastive fine-tuning and hard-negative mining. Training with in-batch negatives plus mined hard negatives (lexically/semantically close but non-answering) is standard in dense retrieval, e.g. DPR’s BM25-mined negatives (Karpukhin et al., 2020); we use TF-IDF at tile granularity instead. Training pairs are generated synthetically by prompting a VLM per tile, in the spirit of InPars (Bonifacio et al., 2022) and Promptagator (Dai et al., 2022) adapted to image tiles. We fine-tune with LoRA (Hu et al., 2021) on both the decoder and the vision tower.

3 Method

A page is rendered once with a headless browser (Appendix D, Figure 2 diagrams the flow); its DOM geometry drives an adaptive tiling step (§3.1); tiles are organized into a two-level graph, optionally spanning several hyperlinked pages (§3.2); a VLM reader is contrastively fine-tuned to score (question, tile) pairs (§3.3); and, at query time, an online controller walks the graph under a budget to decide which nodes the reader actually opens (§3.4).

3.1 DOM-Guided Adaptive Tiling

Rendering and extraction. We load the page in headless Chromium at a fixed viewport width (2048 px), remove interface chrome via a CSS selector list (navigation, footers, Wikipedia furniture), and take a full-page screenshot. For a configurable tag set (h1-h4, p, table, figure, img, ul/ol, blockquote, pre) we read each element’s absolute bounding box directly from the rendered DOM. For non-text containers this is `getBoundingClientRect`; for text-flow tags, whose CSS box spans the full container width by default regardless of actual content, we instead union the client rects of every descendant *text node* (never the element itself, which would also pick up a nested block child’s box). Without this, a paragraph or reference list sharing a row with a floated infobox reports a bounding box running through the infobox’s column even though its text stops well short of it. Hyperlinks are collected from `<p>` elements only, seeding the cross-page edges of §3.2.

Overlap resolution. Nested tags (an `<img>` in a `<figure>`) yield duplicate bounding boxes for the same content; we sort by DOM depth and drop any element already fully contained (1 px tolerance) in a shallower kept element. Containment is not the only overlap source: a floated sidebar and the text sharing its row can genuinely intersect even after the text-flow tightening above. Rather than tile such elements independently and duplicate the shared pixels, we connect any two elements whose bounding-box intersection exceeds a small fraction of the smaller one’s area, and tile each connected component as one crop spanning their union.

Adaptive tiling and patching. Elements are processed top-to-bottom under three rules: (1)

heading pairing – a heading is never a standalone tile, but held and merged with the next element (e.g. h2+p); (2) *merging* – non-heading elements under 80 px are buffered until their union spans that minimum, then flushed as one merged tile; (3) *patching* – an element taller than 2048 px is sliced into $\lceil h/h_{\max} \rceil$ equal-height sub-tiles. This directly targets PixelRAG’s failure mode: a table that would straddle a fixed grid cell is instead one tile (or content-aligned sub-tiles if unusually tall), and a short caption merges with its neighbor rather than orphaning. Each tile keeps a pointer to its source element(s) (tag, text, links), used both for the graph (§3.2) and as ground truth for negative-mining (§3.3). Figure 3 (Appendix D) shows tiles overlaid on a real page.

3.2 Multigranular Graph Construction

Tiles are organized into a two-level directed graph (one page node, one element node per tile) with all five of MAGE-RAG’s relation types, plus one of our own for cross-page connectivity:

- **contains** (page $\rightarrow$ element): every element belongs to its page.
- **reading_order** (element_{*i*} $\rightarrow$ element_{*i+1*}): a spatial heuristic, *not* DOM order – elements are grouped into visual rows (20 px vertical-overlap tolerance), each sorted left-to-right, rows concatenated top-to-bottom, since DOM and visual order diverge under CSS positioning (a sidebar that appears late in the DOM but renders beside the lead paragraph).
- **layout_adjacency** (element $\leftrightarrow$ element): reciprocal edges to the 2D neighbors reading order flattens away (left/right within a row; above/below to the nearest horizontally-overlapping tile in the next row), keeping a multi-column layout’s true neighbors explicit.
- **section_hierarchy** (heading $\rightarrow$ child): each heading points to its nested sub-heading/content, respecting level nesting, forming a section tree distinct from the flat contains relation.
- **semantic_neighbor** (element $\leftrightarrow$ element): TF-IDF-cosine edges between lexically close elements regardless of position, added greedily under a per-tile degree cap. *Optional* (`use_semantic_similarity=True`,

285	off by default) – meant to be toggled to	state=...>text</element> per node)	335
286	compare the controller with and without it	the reader’s structured input, in the spirit of	336
287	(§5), not assumed to help unconditionally,	MAGE-RAG’s structured-rendering step (exam-	337
288	and needing no learned embedder to test as	ple in Appendix D, Figure 6).	338
289	an ablation.		
290	• links_to (element $\rightarrow$ external_page):	3.3 Contrastive Reader Fine-Tuning	339
291	hyperlinks from the first $k=3$ link-bearing	To turn a general-purpose VLM into a tile-level	340
292	text tiles (reading order); beyond that budget,	reader that also scores how well a tile answers a	341
293	links are dropped to bound the corpus crawl’s	query, we fine-tune it contrastively on synthetic	342
294	branching factor.	data, in three stages.	343
295	Every element node starts with a <code>state</code> attribute		
296	(INACTIVE), which the controller (§3.4) evolves at	Synthetic QA generation. For every tile, we	344
297	query time. Figure 4 (Appendix D) shows the re-	prompt the VLM with the tile image alone for one	345
298	sulting graph for the page of Figure 3.	self-contained, factual question answerable <i>only</i>	346
299		from that tile, or <code>SKIP</code> if it carries no exploitable	347
300	Oversized-page splitting. One article’s screen-	content – the recipe of synthetic query generation	348
301	shot exceeded 46,000 px and produced dozens of	for retriever training (InPars, Promptagator; Boni-	349
302	tiles whose sequential QA generation (§3.3) ran	facio et al., 2022; Dai et al., 2022) applied to im-	350
303	for hours with no partial progress saved on fail-	age tiles. Every tile crop is persisted regardless	351
304	ure, since the whole page was one unit of work.	of outcome, so <code>SKIPPED</code> tiles remain available as	352
305	Rather than exclude such pages, one whose ele-	candidate hard negatives.	353
306	ment height exceeds $10\times$ the max tile height is		
307	split into sections <i>before</i> tiling, at <code>h1/h2</code> bound-	Hard-negative mining. For each question, we	354
308	aries (content before the first heading forms a	rank every other tile from the <i>same page</i> by TF-	355
309	preamble section), with any still-oversized section	IDF cosine similarity and keep the top $n=2$ clear-	356
310	further re-chunked by a height budget. Each sec-	ing two false-negative filters: (a) a free lexical fil-	357
311	tion is tiled and graphed as a standalone page	ter rejecting any tile whose text already contains	358
312	would be, and the per-section graphs are merged	the answer verbatim, and (b), optionally, a VLM-	359
313	by node-namespacing, with consecutive sections’	judge filter showing the candidate image to the	360
314	page nodes linked by a new continues rela-	reader VLM and asking whether it can recover the	361
315	tion – distinct from <code>links_to</code> : sequential parts	answer there – catching non-lexical false negatives	362
316	of the <i>same</i> document, not a hyperlink elsewhere.	(paraphrase, logo, chart color) at the cost of one	363
317	Every page’s graph, split or not, exposes an <i>entry</i>	extra VLM call per candidate, so it is opt-in.	364
318	<i>point</i> (the sole <i>page</i> node, or the first section’s)		
319	for the corpus graph below to rewire onto; the con-	Contrastive LoRA fine-tuning. Qwen2-VL is	365
320	troller does not yet cross <code>continues</code> edges auto-	generative, not a similarity encoder; we turn it	366
321	matically (§6). Figure 5 (Appendix D) shows an	into one via PromptEOL (Jiang et al., 2024) (a	367
322	example split into 16 sections.	short fixed instruction, “summarize the above in	368
323		one word,” after the question or a tile image, tak-	369
324	Corpus graph. From a seed page we crawl up	ing the last token’s final-layer hidden state as em-	370
325	to m pages reachable via its <code>links_to</code> edges	bedding), originally proposed for text-only LLMs	371
326	(one hop, non-recursive), build each linked page’s	and applied here to a VLM’s image+instruction in-	372
327	graph independently, and merge by namespacing	put. LoRA ($r=16$, $\alpha=32$, dropout 0.05) adapts	373
328	every node with a URL-derived slug, rewiring	both the decoder’s attention projections and the vi-	374
329	each placeholder external-page node onto the	sion tower’s, since the task requires refining what	375
330	crawled page’s entry point. <code>links_to</code> is the only	the model attends to <i>in the image</i> jointly with how	376
331	cross- <i>document</i> connective tissue, <code>continues</code>	it reads <i>the question</i> . Given B questions each with	377
332	the only cross- <i>section</i> one – the substrate the ev-	one positive and K mined negatives, we minimize	378
333	idence controller (§3.4) operates over.	a temperature-scaled ($\tau=0.05$) InfoNCE loss over	379
334		the $B(1+K)$ tiles in the batch, so each positive	380
	Structured rendering. Any (sub-)graph serial-	competes against its own negatives and, in-batch,	381
	izes its page node and directly-contained elements	every other question’s tiles.	382
	to XML (one <code><element id=... tag=...</code>		




Figure 1: The controller’s state machine. Opening a node consumes budget and may activate unexplored `reading_order` neighbors; pruning is free and applied to any node still ACTIVE once budget is exhausted.

Resolution as an efficiency lever. PixelRAG reports up to $3\times$ token-cost reduction at lower resolution without specifying the filter or resolutions tested. We build the infrastructure to test this directly: every tile can be mirrored via Lanczos resampling at a configurable scale (0.5 default). Since `tiles_manifest.jsonl` and the mined examples store only paths relative to an images root, never pixel dimensions, the same questions/positives/negatives are reused unmodified at both resolutions – comparing them is a matter of pointing fine-tuning at one root or the other, isolating resolution as the only variable (§5).

3.4 Online Evidence Controller

Given a query, the controller decides which element nodes the reader actually reads, rather than the whole page or a static top- k . It evolves each node’s state through the four-state machine of Figure 1, generalizing MAGE-RAG’s Algorithm 1.

Relevance scoring. Every element node is scored against the query by TF-IDF cosine similarity (same mechanism as hard-negative mining, with a pure-Python fallback if `scikit-learn` is unavailable) – a deliberate placeholder for the fine-tuned reader’s own learned similarity (§6).

Policy. Algorithm 1: (1) seed the frontier with the k highest-scoring inactive nodes; (2) repeatedly take the highest-scoring *active* node, pruning it for free below an opening threshold or opening it (one unit of budget B) and activating still-inactive `reading_order` neighbors that clear an activation threshold; (3) once budget is exhausted or no active node remains, prune whatever is left active. Neighbors are only considered *after* their node is opened, so the frontier grows along reading order conditioned on what has already proven relevant,

Algorithm 1 Budgeted best-first evidence control

Require: graph G , query q , budget B , seed count k , thresholds $\theta_{\text{open}}, \theta_{\text{act}}$

- 1: score every element node by $\text{sim}(q, \text{text}(n))$
- 2: activate the k highest-scoring inactive nodes $\triangleright$ seed
- 3: $b \leftarrow 0$
- 4: **while** $b < B$ and some node is ACTIVE **do**
- 5: $n \leftarrow \arg \max \text{score over ACTIVE nodes}$
- 6: **if** $\text{score}(n) < \theta_{\text{open}}$ **then**
- 7: $n \leftarrow \text{PRUNED}$ $\triangleright$ free
- 8: **else**
- 9: $n \leftarrow \text{OPENED}; b \leftarrow b + 1$; record its text as evidence
- 10: **for each** inactive `reading_order` neighbor n' of n **do**
- 11: **if** $\text{score}(n') \geq \theta_{\text{act}}$ **then**
- 12: $n' \leftarrow \text{ACTIVE}$
- 13: **end if**
- 14: **end for**
- 15: **end if**
- 16: **end while**
- 17: prune every node still ACTIVE $\triangleright$ finalize
- 18: **return** concatenation of OPENED nodes’ text, in reading order

rather than a fixed global ranking – what distinguishes it from static top- k retrieval over the same scores (ablation in §5).

Output. The concatenation of OPENED nodes’ text, in reading order, is the final evidence passed to the downstream reader, with the controller’s action log available for inspection.

4 Implementation

The pipeline is Python, one module per stage of §3: rendering/extraction via Playwright (`dom_extraction.py`); tiling and pruning over Pillow crops (`tiling.py`); graph construction with `networkx` (`graph_builder.py`); the reader VLM via Ollama (`qwen3-vl`, default, no GPU needed) or Hugging Face transformers (`Qwen2-VL-7B-Instruct`) for LoRA fine-tuning, every call cached to disk by a hash of (image, prompt, decoding-config) – safe since decoding is greedy (temperature 0), so a retried page never re-pays a call already made (`vlm_client.py`); synthetic QA generation (`qa_generation.py`); hard-negative mining via `scikit-learn`

TF-IDF with a dependency-free fallback (`hard_negative_mining.py`); and contrastive LoRA fine-tuning with `peft/torch`, target modules discovered by introspecting the loaded model rather than hard-coded, so the code does not silently no-op if `transformers`’ internal naming changes (`lora_finetune.py`).

Oversized-page splitting, composition, and concurrency. `sectioning.py` implements the h1/h2 split as a pure function over already-extracted elements, unit-testable without Playwright. `corpus_builder.py` separates network-bound fetch from a synchronous core that tiles, decides whether to split, and composes per-section graphs, exposing an `entry_page` on every graph; multi-page (`links_to`) and multi-section (`continues`) composition share the same node-namespacing machinery. QA generation issues per-tile VLM calls concurrently within a page (a thread pool sized to Ollama’s default parallelism); the controller does the same on the query path, since which nodes open depends only on relevance scores computed upfront, never on a VLM response, so their reads have no inter-dependency once that set is fixed.

Batch construction and resolution ablation. `scripts/build_dataset.py` drives the pipeline over a page list with retrying fetches; its unit of work is a *page unit* (a page, or one split section), checkpointed independently to `contrastive_examples.jsonl` – a failure partway through an oversized page costs only the section in flight, motivated by an early run that lost hours of work when the whole page was the retry unit. `image_compression.py` mirrors tile images at a reduced scale under a separate root with identical filenames and relative paths, so switching a training run between resolutions (§3.3) is a single `--images-root` flag.

Tiling ablation and multi-hop data. `scripts/grid_baseline.py` reimplements the blind grid as a module independent of `tiling.py`; `scripts/dom_vs_grid_experiment.py` re-renders each page, builds both tilings, and ranks existing questions against both by TF-IDF (§5). `src/multihop_qa.py` builds cross-page examples from HotpotQA (Yang et al., 2018) records streamed via the `datasets` library (only sampled questions are fetched);

`scripts/build_multihop_dataset.py` reports unreachable-article vs. unmatched-sentence drops separately. Positive-tile attribution uses stopword-filtered token overlap rather than exact substring matching, since the latter rarely survives a decade of article edits since HotpotQA’s 2017 source dump.

Demonstration and tests. A notebook runs every stage end-to-end with interactive `pyvis` graph visualizations. Dependency-light unit tests cover the geometric pipeline, graph construction and splitting, hard-negative mining (lexical and VLM-judge, stub client), the VLM cache and call concurrency (asserting calls actually overlap in time), and resolution compression, all on synthetic fixtures needing no network, VLM, or GPU. Playwright-backed tests check the text-flow bounding-box measurement (float-ed sidebar fixture) and the full geometric pipeline end-to-end (DOM extraction through XML export) against a bundled local HTML fixture – neither needs a paid or rate-limited service.

5 Evaluation Protocol and Status

This section reports status honestly: the pipeline of §3 runs end-to-end on live pages, but the trained reader and full benchmark suite are not yet in place. We report one preliminary lexical ablation now rather than fill this section with numbers ahead of the actual runs.

Benchmarks. MAGE-RAG’s protocol evaluates on LongDocURL (Deng et al., 2025) and MMLongBench-Doc (Ma et al., 2024b) (Table 1), but both are PDF benchmarks, and every input our method needs (DOM bounding boxes, CSS-selector chrome removal, h1/h2 splitting) presupposes a renderable DOM a PDF does not have – reporting a number on either without resolving that mismatch would not be a controlled comparison. We instead plan an HTML-native source: WebSRC (Chen et al., 2021) (human-annotated HTML + screenshots) or HotpotQA (Yang et al., 2018), whose `supporting_facts` name the exact sentences, in two articles, needed per question – multi-hop by construction and the most direct fix for our own synthetic QA’s single-hop bias (§6). We have implemented the HotpotQA path (`src/multihop_qa.py`, §6): 24/40 sampled questions yielded a complete cross-page example; scaling it is next. We do not evaluate on SimpleQA (an earlier draft’s choice): with no supplied document, it cannot exercise tiling, the graph, or the

Benchmark	Metric	MAGE-RAG
LongDocURL	Accuracy	52.75
MMLongBench-Doc	Acc. / F1	53.26 / 51.19

Table 1: Published MAGE-RAG (Zuo et al., 2026) numbers, the reference our controller (§3.4) generalizes. Our results on the same benchmarks are ongoing work.

	<i>n</i>	R@1	MRR	MRR/chance
<i>All attributable pairs</i>				
DOM	223	0.543	0.666	14.2×
Grid	115	0.730	0.824	3.7×
<i>Paired (109 pairs attributable in both)</i>				
DOM	109	0.505	0.661	15.8×
Grid	109	0.725	0.821	3.9×

Table 2: DOM-guided vs. blind grid, same TF-IDF scorer, 18 pages / 425 questions. MRR/chance = MRR ÷ mean chance-level MRR ($1/n_{\text{items}}$) of that row’s own pool, since grid’s pool is $\sim 5\times$ smaller (Appendix C) and inflates raw Recall/MRR mechanically.

controller.

Ablation (a): DOM-guided tiling vs. a blind fixed-height grid. The lexical scorer already used for hard-negative mining and controller relevance (§3.3, §3.4) lets us run this ablation without waiting on the trained reader. `scripts/grid_baseline.py` reimplements a PixelRAG-style blind grid (1,024px bands, PixelRAG’s own reported tile height, at our render width); `scripts/dom_vs_grid_experiment.py` re-renders 18 single-unit pages (the 6 split pages need `sectioning.py` re-run too and are left for a follow-up), tiles each both ways, and ranks every existing `qa_pairs.jsonl` question against both tile sets by TF-IDF cosine – same scorer, only the tiling differs. Both conditions are restricted to the region covered by the first 25 DOM tiles (matching how the QA pairs were generated); a question is scored in a condition only if its answer string occurs in one of that condition’s tiles, else dropped from that condition rather than counted as a miss.

Raw numbers say grid is more accurate – the opposite of this paper’s claim – but are confounded: grid’s pool is $\sim 5\times$ smaller (Appendix C), which improves Recall/MRR mechanically regardless of ranking quality. Normalizing by each row’s own chance-level MRR reverses the conclusion: DOM-guided tiling is **14–16×** better than chance on its

Statistic	Value
Distinct Wikipedia articles	19
Page units (articles + split sections)	93
Tiles extracted	1,103
QA pairs generated	970
Tiles with no usable QA (SKIP)	133 (12.1%)
Contrastive examples (QA + hard negatives)	534

Table 3: Contrastive dataset size (§3.3), computed from `tiles_manifest.jsonl`, `qa_pairs.jsonl`, `contrastive_examples.jsonl` (lexical filter only, no VLM judge). 6 of 19 articles were split into sections (§3.2).

own pool, against grid’s **3.7–3.9×** on its own – a gap that survives on the paired subset too, so it is not an artifact of differing question sets. We read this as a real retrieval advantage from content-aligned boundaries, larger than raw Recall@*k* suggests, though a preliminary signal (single lexical scorer, self-generated single-hop questions, §6) rather than a final result. Ablations (b)/(c) still require the trained reader.

Status. All stages (§3) are implemented, tested, and demonstrated end-to-end on live pages (§4). QA-pair generation has scaled to the multi-page corpus of Table 3. Scaling surfaced a real failure mode – our local reader occasionally reasoned until it exhausted Ollama’s context window, silently emitting an empty, discarded response, fixed by widening the window and disabling reasoning for this step – and a real cost: ~ 2 minutes/tile on local (Apple Silicon) hardware, now partly mitigated by concurrent VLM calls (§4); Appendix B turns this into a compute-budget estimate for fine-tuning itself. Full results, including ablations (b)/(c), follow once both move to rented GPU compute.

6 Limitations

Relation coverage is complete, consumption is not. All five of MAGE-RAG’s relation types are instantiated (§3.2), but the controller’s frontier expansion (Algorithm 1) only ever follows `reading_order`: `layout_adjacency`, `section_hierarchy`, and `continues` are built and visualized but never queried, and `semantic_neighbor` is off by default pending an ablation. Operationally the graph the controller walks is closer to a chain than the structure it is built from. Type-aware, per-relation expansion is the natural next step, and a higher-value one than adding further relations.

Single-hop self-generated data cannot exercise the graph – addressed via HotpotQA, at modest scale. §3.3’s QA generation asks for a question answerable *only* from the tile shown, by construction, so an oracle controller could answer every such question from one node: the graph cannot be shown to help on this data. `src/multihop_qa.py` closes this gap with HotpotQA (Yang et al., 2018) instead: each question’s `supporting_facts` names two Wikipedia articles and the sentences in each needed to answer it, so we render both, keep tiles overlapping a supporting sentence (token overlap, tolerant of rewording since HotpotQA’s 2017 source dump) as positives, and mine hard negatives by TF-IDF pooled across both articles’ tiles – unavoidably cross-page, since the positives themselves span two pages. On 40 sampled questions, 24 (60%) yielded a complete example, every rejection an unmatched sentence rather than a missing article. Scaling this and using it in training/evaluation remains future work.

Scoped to semantically marked-up pages. Chrome removal uses a MediaWiki-specific selector list, and every result here is on Wikipedia; we do not claim this works unmodified on JS-hydrated SPAs or `div-soup` layouts, the regime web-agent work such as WebArena (Zhou et al., 2023), Mind2Web (Deng et al., 2023), and See-Act (Zheng et al., 2024) already studies from the acting side. The open web is future work.

Same-model-family generation, judging, and reading. Synthetic questions, the optional false-negative judge, and (a different generation of) the fine-tuned reader all come from the `qwen3-vl/Qwen2-VL` family, so any measured gain could partly reflect the generator’s idiosyncrasies rather than genuine document understanding; hand-verifying generated pairs, and eventually human-authored questions, is needed to rule this out.

Arbitrary slicing, lexical scoring, and other gaps. Patching (§3.1) splits an oversized element into equal-height sub-tiles regardless of content – a narrower version of PixelRAG’s own failure mode; a dynamic, resolution-preserving scheme (Qwen2-VL’s AnyRes) is the natural fix. The controller’s TF-IDF relevance scoring (§3.4) misses paraphrases and cross-modal cues the fine-tuned reader would catch; wiring in the reader’s

own similarity is the natural next iteration, unifying mining and control. The VLM-judge false-negative filter is opt-in and unused so far; hard-negative mining is single-page outside the HotpotQA path above; the corpus crawl is one hop, non-recursive; and all hyperparameters (Appendix A) are centralized but chosen once, not swept, with pixel-valued ones tied to a single 2048 px viewport.

7 Conclusion

DOM-Graph RAG grounds two ideas in visual and graph-structured RAG – PixelRAG’s screenshot tiles and MAGE-RAG’s budgeted graph controller – in the DOM a rendered web page already provides for free: content-aligned tiling replaces a blind pixel grid, a multigranular graph replaces a flat tile list, a contrastively fine-tuned reader replaces a generic VLM, and a four-state evidence controller replaces reading the whole page or a static top-*k*. Every stage runs end-to-end on live pages, and a first lexical ablation (§5) suggests the tiling carries a real retrieval advantage once pool-size effects are controlled for. Next: an HTML-native benchmark, the remaining ablations, and closing the gaps of §6 – most importantly, the controller’s lexical scoring replaced with the fine-tuned reader’s own learned similarity.

Ethics Statement

The corpus builder fetches a seed Wikipedia page plus a small, bounded number of linked pages, with retry-and-backoff rather than aggressive re-fetching; we crawled only Wikipedia, built for programmatic access at this volume, though a production deployment should prefer the Wikimedia REST API or a static dump over live scraping, which this prototype does not do by default. Wikipedia text is CC BY-SA; embedded images carry heterogeneous, some non-free, licenses, so the screenshots/tiles produced are derivative works we do not release here (a future release would need per-image licensing). Synthetic questions (§3.3) are VLM-produced and unverified, and should be labeled as such if released (§6). This work involves no human subjects or user study.

References

Luiz Bonifacio, Hugo Abonizio, Marzieh Fadaee, and Rodrigo Nogueira. 2022. [Inpars: Data augmenta-](#)

708	tion for information retrieval using large language	the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP).	764
709	models. <i>Preprint</i> , arXiv:2202.05144.		765
710	Xingyu Chen, Zihan Zhao, Lu Chen, Danyang Zhang,	Junlong Li, Yiheng Xu, Lei Cui, and Furu Wei. 2022.	766
711	Jiabao Ji, Ao Luo, Yuxuan Xiong, and Kai Yu. 2021.	MarkupLM: Pre-training of text and markup lan-	767
712	WebSRC: A dataset for web-based structural read-	guage for visually-rich document understanding. In	768
713	ing comprehension. In <i>Proceedings of the 2021 Con-</i>	<i>Proceedings of the 60th Annual Meeting of the Asso-</i>	769
714	<i>ference on Empirical Methods in Natural Language</i>	<i>ciation for Computational Linguistics (ACL)</i> .	770
715	<i>Processing (EMNLP)</i> .		
716	Jaemin Cho, Debanjan Mahata, Ozan Irsoy, Yu-	Xueguang Ma, Sheng-Chieh Lin, Minghan Li, Wenhu	771
717	jie He, and Mohit Bansal. 2024. M3docrag:	Chen, and Jimmy Lin. 2024a. Unifying multimodal	772
718	Multi-modal retrieval is what you need for multi-	retrieval via document screenshot embedding. In	773
719	page multi-document understanding. <i>Preprint</i> ,	<i>Proceedings of the 2024 Conference on Empirical</i>	774
720	arXiv:2411.04952.	<i>Methods in Natural Language Processing (EMNLP)</i> .	775
721	Zhuyun Dai, Vincent Y. Zhao, Ji Ma, Yi Luan, Jianmo	Yubo Ma, Yuhang Zang, Liangyu Chen, Meiqi Chen,	776
722	Ni, Jing Lu, Anton Bakalov, Kelvin Guu, Keith B.	Yizhu Jiao, Xinze Li, Xinyuan Lu, Ziyu Liu, Yan	777
723	Hall, and Ming-Wei Chang. 2022. Promptagator:	Ma, Xiaoyi Dong, Pan Zhang, Liangming Pan, Yu-	778
724	Few-shot dense retrieval from 8 examples. <i>Preprint</i> ,	Gang Jiang, Jiaqi Wang, Yixin Cao, and Aixin Sun.	779
725	arXiv:2209.11755.	2024b. Mmlongbench-doc: Benchmarking long-	780
		context document understanding with visualizations.	781
		<i>Preprint</i> , arXiv:2407.01523.	782
726	Chao Deng, Jiale Yuan, Pi Bu, Peijie Wang, Zhong-	Yichuan Wang, Zhifei Li, Zirui Wang, Paul Teiletche,	783
727	Zhi Li, Jian Xu, Xiao-Hui Li, Yuan Gao, Jun	Lesheng Jin, Matei Zaharia, Joseph E. Gonzal-	784
728	Song, Bo Zheng, and Cheng-Lin Liu. 2025. Long-	ez, and Sewon Min. 2026. Pixelrag: Web	785
729	DocURL: a comprehensive multimodal long docu-	screenshots beat text for retrieval-augmented gen-	786
730	ment benchmark integrating understanding, reason-	eration. <i>Preprint</i> , arXiv:2606.28344. UC Berkeley /	787
731	ing, and locating. In <i>Proceedings of the 63rd Annual</i>	BAIR / Berkeley NLP. Code: github.com/StarTrail-	788
732	<i>Meeting of the Association for Computational Lin-</i>	org/PixelRAG.	789
733	<i>guistics (Volume 1: Long Papers)</i> , pages 1135–1159,		
734	Vienna, Austria. Association for Computational Lin-	Yiheng Xu, Minghao Li, Lei Cui, Shaohan Huang,	790
735	guistics.	Furu Wei, and Ming Zhou. 2020. LayoutLM: Pre-	791
736	Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen,	training of text and layout for document image	792
737	Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su.	understanding. In <i>Proceedings of the 26th ACM</i>	793
738	2023. Mind2web: Towards a generalist agent for the	<i>SIGKDD International Conference on Knowledge</i>	794
739	web. <i>Preprint</i> , arXiv:2306.06070. NeurIPS 2023.	<i>Discovery & Data Mining (KDD)</i> .	795
740	Darren Edge, Ha Trinh, Newman Cheng, Joshua	Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Ben-	796
741	Bradley, Alex Chao, Apurva Mody, Steven Truitt,	gio, William W. Cohen, Ruslan Salakhutdinov, and	797
742	and Jonathan Larson. 2024. From local to global: A	Christopher D. Manning. 2018. HotpotQA: A	798
743	graph RAG approach to query-focused summariza-	dataset for diverse, explainable multi-hop question	799
744	tion. <i>Preprint</i> , arXiv:2404.16130. Microsoft.	answering. In <i>Proceedings of the 2018 Conference</i>	800
745	Manuel Faysse, Hugues Sibille, Tony Wu, Bilel Om-	<i>on Empirical Methods in Natural Language Process-</i>	801
746	rani, Gautier Viaud, Céline Hudelot, and Pierre	<i>ing (EMNLP)</i> .	802
747	Colombo. 2024. Colpali: Efficient document re-		
748	trieval with vision language models. <i>Preprint</i> ,	Shi Yu, Chaoyue Tang, Bokai Xu, Junbo Cui, Jun-	803
749	arXiv:2407.01449.	hao Ran, Yukun Yan, Zhenghao Liu, Shuo Wang,	804
750	Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan	Xu Han, Zhiyuan Liu, and Maosong Sun. 2024.	805
751	Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,	Visrag: Vision-based retrieval-augmented gen-	806
752	and Weizhu Chen. 2021. Lora: Low-rank	eration on multi-modality documents. <i>Preprint</i> ,	807
753	adaptation of large language models. <i>Preprint</i> ,	arXiv:2410.10594.	808
754	arXiv:2106.09685.		
755	Ting Jiang, Shaohan Huang, Zhongzhi Luan, Deqing	Boyuan Zheng, Boyu Gou, Jihyung Kil, Huan Sun,	809
756	Wang, and Fuzhen Zhuang. 2024. Scaling sentence	and Yu Su. 2024. GPT-4V(ision) is a generalist	810
757	embeddings with large language models. In <i>Find-</i>	web agent, if grounded. <i>Preprint</i> , arXiv:2401.01614.	811
758	<i>ings of the Association for Computational Linguis-</i>	ICML 2024.	812
759	<i>tics: EMNLP 2024</i> , pages 3182–3196.		
760	Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick	Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou,	813
761	Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and	Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue	814
762	Wen-tau Yih. 2020. Dense passage retrieval for	Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Gra-	815
763	open-domain question answering. In <i>Proceedings of</i>	ham Neubig. 2023. Webarena: A realistic web envi-	816
		ronment for building autonomous agents. <i>Preprint</i> ,	817
		arXiv:2307.13854. ICLR 2024.	818

Yilong Zuo, Xunkai Li, Jing Yuan, Qiangqiang Dai, Hongchao Qin, and Ronghua Li. 2026. [MAGE-RAG: Multigranular adaptive graph evidence for agentic multimodal RAG in long-document QA](#). *Preprint*, arXiv:2606.15906. Code: github.com/laonuo2004/MAGE-RAG.

A Hyperparameters

Parameter	Value
Viewport / max tile width	2048 px
Max tile height (patching)	2048 px
Min tile height (merging)	80 px
Row-grouping tolerance	20 px
Nesting tolerance	1 px
Overlap-grouping threshold	15% of smaller area
Link-bearing tiles kept (k)	3
Oversize-split threshold	$10 \times$ max tile height
Mined hard negatives (n)	2
Semantic-neighbor threshold	0.2 cosine
Semantic neighbors / tile (max)	3
Downscale factor (ablation)	0.5
LoRA r / α / dropout	16 / 32 / 0.05
InfoNCE temperature τ	0.05

Table 4: All pipeline hyperparameters (§3.1–§3.3). Each is a single documented, centralized value, not swept; see §6.

B Compute budget for contrastive fine-tuning

Estimated from the training loop’s structure (§3.3), not measured: at batch size 4, $n=2$ hard negatives, 3 epochs over the 534 examples of Table 3, one optimizer step embeds 4 questions and 12 tile images; the current implementation calls the model once per question and once per image rather than as a single batched forward pass, so one step costs 16 sequential forward passes through Qwen2-VL-7B-Instruct, and a full run costs $402 \text{ steps} \times 16 = 6,432$ forward passes, all contributing to one shared backward per step. Weights alone are ≈ 14 GB in `bfloat16`; LoRA keeps the optimizer state small, but activation memory scales with the frozen base model’s forward/backward graph across all 16 unbatched calls held simultaneously before the shared backward. We expect a single 40–80 GB datacenter GPU (A100/H100 class) to suffice at this dataset size – no multi-GPU training is warranted – sized conservatively until peak memory is profiled, since the vision encoder can still emit a large token count for our largest tiles (2048×2048 px) before its own internal resizing caps it. At an assumed 150–400 ms per forward call, the 6,432

calls total 16–43 minutes of forward compute alone; including backward and the overhead of 16 unbatched calls per step, we would not be surprised by anything from under an hour to several hours wall-clock for the full run. Batching `embed_questions/embed_tiles` into one padded forward pass per step is the highest-value change before renting GPU time, reducing both the estimate and its uncertainty; we flag it as the next engineering step rather than fold an unbatched-loop number into a claimed result.

C Tiling ablation: cost detail

	Tiles/page	Px/tile	Total px/page
DOM	24.0	158,072	3,793,738
Grid	4.9	2,081,117	10,174,350

Table 5: Tiling cost, same 18 pages as Table 2. DOM-guided tiling produces more, smaller tiles, yet a smaller *total* pixel budget per page ($2.7\times$ less) – a direct consequence of the text-flow bounding-box measurement of §3.1, which crops text tiles to the width their content actually occupies instead of the full page width every grid band uses regardless of content. If image tokens scale with pixel area, as they do for Qwen2-VL’s vision encoder, DOM-guided tiling need not trade accuracy for token cost against a blind grid.

Attach rate (the fraction of questions whose answer string survives a tiling’s 200-character text preview) is 52.5% for DOM (223/425) against 27.1% for grid (115/425) – consistent with grid bands aggregating far more raw content into the same fixed text budget, diluting the specific snippet a lexical scorer needs to surface.

D Additional figures




Figure 2: Pipeline overview. Rendering and tiling happen once per page; the evidence controller runs once per query, deciding which graph nodes the reader opens.




Figure 3: DOM-guided tiles (colored by content category, numbered in reading order) overlaid on the rendered page for `en.wikipedia.org/wiki/George_P._Gunn`. Tile boundaries follow the DOM’s own element boundaries: the infobox table (tile 2) is never split mid-row and stays disjoint from the body text and reference list that visually share its column range (tiles 1, 3–6), each section heading is paired with its first content block – a paragraph for tiles 3–5, the reference list itself for tile 6 – rather than emitted alone, and the trailing category footer (tile 7) is merged into its own tile rather than left as a flood of near-empty fragments.




Figure 4: Page/element graph for the same page as Figure 3, laid out by real horizontal position (x) and reading-order rank (y). `reading_order` (solid black) and `contains` (light gray) are the two relations MAGE-RAG also has; `layout_adjacency` (dotted blue) and `section_hierarchy` (dashed red) are the two we add. Most headings here are paired with their content into one tile (§3.1), so `section_hierarchy` shows only where trailing non-heading content (the category footer) attaches to the last open section, rather than between the (sibling, not nested) section headings themselves.




Figure 5: Page-level view of the graph for `en.wikipedia.org/wiki/The_Beatles` (278 DOM elements, split into 16 sections at `h1/h2` boundaries – only the first 6 are shown for legibility). Each diamond is a page node, one per section; `continues` edges (black) chain them in document order; `links_to` edges (pink) reach external pages, collapsed here from their true element-level source (§3.2) to stay at page granularity. The leftmost node is the entry point.


```
<page url="..." title="...">
  <element id="tile_0000" tag="merged"
    state="inactive">...</element>
  <element id="tile_0003" tag="p"
    state="opened">Gunn was born
    on October 11, 1903...</element>
  <element id="tile_0004" tag="h2+p"
    state="inactive">...</element>
  ...
</page>
```

Figure 6: Excerpt of the XML rendering of a page graph after the evidence controller has run on the query “When and where was George Purnell Gunn born?”. Only the paragraph tile carrying the birth date/place has been opened; the rest remains inactive or pruned.